



CMPExplore User guide

Version 0.2.1, 04 Feb 2014

Contents

I.	Introduction	2
II.	Build and Installation	2
III.	A Quick Starter: Running an Exploration	3
IV.	Configuring the Design Space	4
V.	Workload Specification.....	6
VI.	Command Line Arguments	7
VII.	Exploration Mode.....	8
VIII.	Modeling a Particular Configuration	9
IX.	Conversion Mode	11
X.	Hidden Parameters (require recompilation).....	12
XI.	History	13
XII.	Contacts	13

I. Introduction

`CMPexplore` is a framework for architectural exploration of hierarchical chip multiprocessors (CMPs). Given a description of the architectural design space in terms of the **models** of CMP components (e.g. core architectures, cache sizes, interconnect topologies, etc.) and **constraints** (e.g. area and power) the tool searches for the architectures that maximize the system throughput in instructions-per-cycle, IPC.

The assumptions and details of the models behind the implementation can be found in the following papers:

N. Nikitin, J. de San Pedro and J. Cortadella.

"Architectural Exploration of Large-Scale Hierarchical Chip Multiprocessors."

In IEEE TCAD, Oct 2013.

N. Nikitin, J. de San Pedro, J. Carmona and J. Cortadella.

"Analytical performance modeling of hierarchical interconnect fabrics."

In Proc. The 6th ACM/IEEE International Symposium on Networks-on-Chip (NOCS), May 2012.

The tool currently supports two-level interconnects with a global mesh at the top-level, and local buses, crossbars, uni- and bi-directional rings at the cluster level.

II. Build and Installation

To build the exploration tool just run `make` in the `cmpexplore` directory. You should finally see the `./cmpexplore` executable.

Note: this release includes the **newmat10D** library, available at <http://www.robertnz.net/nm10.htm>

III. A Quick Starter: Running an Exploration

Invoke the following command from the current tool directory:

```
./cmpexplore -print 10 -dump_configs
```

This will run an *exhaustive* exploration of the design space given by the configuration file `./config.txt`. Full description of the configuration file follows in the next section. As exploration has finished, it prints out 10 best architectures (as given by the `-print` argument), ordered by the best throughput. The argument `-dump_configs` will store the architectures in the `test` directory. The output will look as follows:

```
***** Stats: ReportAllConfigs *****
x= 5; y= 4; P1= 6; L1=0.064; L2= 1; P#=120; L3= 60MB; Ic=bRing; A=342.7; Pwr = 179.1 Thr = 82.783
x= 6; y= 5; P1= 4; L1=0.064; L2= 1; P#=120; L3= 60MB; Ic=bRing; A=345.2; Pwr = 191.5 Thr = 82.618
x= 7; y= 4; P1= 4; L1=0.128; L2= 1; P#=112; L3= 56MB; Ic=bRing; A=329.4; Pwr = 184.3 Thr = 79.528
x= 4; y= 4; P1= 7; L1=0.128; L2= 1; P#=112; L3= 64MB; Ic=bRing; A=334.4; Pwr = 168.8 Thr = 79.273
x= 8; y= 7; P1= 2; L1=0.128; L2= 1; P#=112; L3= 56MB; Ic=bRing; A=336.4; Pwr = 213.6 Thr = 78.648
x= 7; y= 4; P1= 4; L1=0.064; L2= 1; P#=112; L3= 56MB; Ic=bRing; A=322.2; Pwr = 181.1 Thr = 77.611
x= 4; y= 4; P1= 7; L1=0.064; L2= 1; P#=112; L3= 80MB; Ic=bRing; A=343.2; Pwr = 166.7 Thr = 77.385
x= 8; y= 7; P1= 2; L1=0.064; L2= 1; P#=112; L3= 56MB; Ic=bRing; A=329.2; Pwr = 210.3 Thr = 76.748
x= 9; y= 6; P1= 2; L1=0.128; L2= 1; P#=108; L3= 54MB; Ic=bRing; A=324.4; Pwr = 207.5 Thr = 76.099
x= 6; y= 6; P1= 3; L1=0.128; L2= 1; P#=108; L3= 72MB; Ic=bRing; A=337.9; Pwr = 189.6 Thr = 75.899
790 configurations explored
Total time = 28.718
```

- The **x** and **y** parameters show the dimensions of global mesh.
- The next three parameters define the core (a.k.a. processor) configuration of one cluster: the name and the number of cores *per cluster* is given for each core type (e.g. P1 = 6), followed by the L1 and L2 cache sizes in MB *per core*.
- **P#** shows the total number of cores per chip.
- **L3** is the total amount of L3 cache distributed over the chip.
- **Ic** is the type of cluster interconnect: **Bus-N/xBar-N** is for a bus/crossbar with the latency of N cycles, **uRing** is for a uni-directional ring and **bRing** is for a bi-directional ring.
- The **A**, **Pwr** and **Thr** columns define the architecture metrics, as obtained by the analytical model, namely the area, power and throughput (IPC).

E.g. the best architecture for this exploration example is a 5x4 mesh with 120 cores, distributed among 20 clusters, each consisting of 6 P1 cores with 64KB L1 and 1MB L2. The total amount of L3 cache is 60Mb. The modeled power and throughput of this architecture are about **179 W** and **83 IPC**, respectively.

IV. Configuring the Design Space

The parameters of the design space for exploration have to be specified in the configuration file, such as `config.txt`, found in the same directory. The description of every parameter follows.

<i>Parameter name</i>	<i>Description</i>	<i>Value specification</i>
gMeshDimX	Global mesh X-dimensions	list of integers
gMeshDimY	Global mesh Y-dimensions	list of integers
clusterIcType	Type of the cluster interconnect	bus, xbar, uring, bring
maxArea	The chip area constraint [mm ²]	Real
processor	Introduces a processor (core) type with a set of parameters, listed below:	string (required keyword)
Name	Core type name	string
Area	Area of the core [mm ²]	real
OoO	In-order (0) vs. out-of-order (1) architecture	integer (0 or 1)
Freq	Core frequency [GHz]	real
Epi	Energy per instruction [nJ]	real
Pleak	Leakage power [W]	real
l1Size	Possible L1 cache sizes	list of reals
l2Size	Possible L2 cache sizes	list of reals
maxProcAreaPerCluster	Maximum fraction of cluster area occupied by the cores (1.0 means all area can be used)	real (in the 0.0-1.0 range)
numL3PerCluster	Number of L3 cache modules per cluster	list of integers
memDensity	Area of 1Mb of the cache memory [mm ²]	Real
wlFile	Workload specification file (see Sect. Workload Specification)	path (string)
l3LatencyOfSize	Dependency of the cache latency on cache size	function (see Note 3 below)
l3ShareDegreeOfPNum	Dependency of the cache sharing degree (number of cores sharing a cache line on average) on the total number of cores on chip	Function
busTimeOfClusterArea	Dependency of the cluster bus latency on the cluster size	Function
xBarDelayOfClusterArea	Dependency of the cluster crossbar latency on the cluster size	Function
meshLinkDelayOfClusterArea	Dependency of the mesh link latency on the cluster size	Function
l3AccessEnergyOfSize	Dependency of the energy of accessing L3 cache on the cache size	Function
l3LeakagePowerOfSize	Dependency of the leakage power of L3 cache on the cache size	function
memCtrlAccessEnergy	Energy of accessing a memory controller [nJ]	real
memCtrlLeakagePower	Leakage power of memory controller [W]	real
frequency	Master clock frequency (memory/IC) [GHz]	real
linkWidth	Interconnect link width	integer (supported: 128, 256, 512, 1024)

Note 1: we currently measure all cache sizes in megabytes and use real values like 0.064, 0.128 for the sizes of 64Kb and 128Kb. This is an approximation done for simplicity of implementation, yet providing readable values, and will be fixed in the future version.

Note 2: to make the parameter specification easier, one can use ranges of **integer** values: notation [m-n] will be internally converted into the sequence of integers {m, m+1, ..., n-1, n}. Range notation can be combined with *enumeration*; hence all the specifications below are *equivalent*:

gMeshDimX = 4, 5, 6, 7, 8, 9, 10

gMeshDimX = [4-10]

gMeshDimX = 4, [5-8], 9, 10

Note 3: function specification. In some cases we need to specify how one parameter depends on another. When a discrete range of values is to be explored, a table-based specification may be enough. But for the purposes of continuous range exploration (as well as for shorter notations), we introduce *functions*. Consider the following definition:

```
l3LatencyOfSize = piecewise (0.064 2) (0.128 3) (0.256 4) (0.512 5) \
                             (1 6) (2 7) (4 8) (8 9) (12 10)
```

Here the dependency of the cache latency on the cache size is defined as a linear piecewise function, based on the set of 9 points. This definition is equivalent to the following table:

Cache size	64K	128K	256K	512K	1M	2M	4M	8M	12M
Cache latency	2	3	4	5	6	7	8	9	10

The piecewise function performs linear interpolation for the values between the specified points, as well as constant extrapolation out of the specified region. For instance, for the 96K the latency will be $2 + (3-2)/2 = 2.5$ cycles (note that it will be further ceiled to 3 cycles in the tool, but this is somewhat independent of the function behavior). For 32K the latency will be 2 cycles, for the 15M the latency will be 10 cycles (constant extrapolation).

Apart from the piecewise function, several other function types are supported:

linear k b – is a linear dependency in the form $k \cdot x + b$

sqrt k – is a square root dependency in the form $k \cdot \sqrt{x}$

powerlaw k a – is a power-law dependency in the form $k \cdot (x^a)$

const k – is a function returning constant value, irrespective of x

V. Workload Specification

Since version 0.2 the tool is able to perform exploration for multiple workloads and select the best architecture according to the *maximum aggregate throughput*. The latter is defined by default as the average throughput which configuration delivers for the set of given workloads. However, if required, the aggregating function can be redefined in the code: `src/stat/StatConfig.cpp`, `StatConfig::AggWIObjective()`.

The `wlFile` parameter in the configuration file is used to point to a file with workload specification. Note that the file has to contain definition of at least one workload application. For every application, the following fields have to be defined, where the workload name always has to be present and go first.

Parameter name	Description	Value specification
wlName	Full name of the workload application	string (required keyword)
wlShortName	Short name (alias)	string
missRatioOfMemSize	Dependency of the miss ratio of executed tasks on the cache size	function (see Note 3 in Sect. Configuring the Design Space)
IPC	Ideal throughput of every type of processor, defined in the configuration file, when executing this application	Comma-separated list of reals
MPI	Fraction of memory references per instruction, characterizing the workload (for compatibility, this parameter also has to be specified for every processor type)	Comma-separated list of reals (in the 0.0-1.0 range)

As an example, consider `wl_spx_nmd_3cpu.txt` found in the project directory. This file defines two workload applications, *soplex* and *namd* from *SPEC CPU2006*. Three processor types are assumed for exploration, such as those described in the configuration file `conf_demo_1.5b.txt`.

VI. Command Line Arguments

Argument	Description
-test <filename>	Model a particular *.cmp configuration given by the filename. See section VII.
-config <filename>	Specify configuration of the design space in filename. Default value = ./config.txt
-exp_mode <string>	Exploration mode: exhaustive ("ex"), simulated annealing ("sa") or extremal optimization ("eo"). Default value = ex.
-max_power <real>	Maximum power constraint. Configuration whose power exceeds max_power will not appear in the final list. Default value = 10 ⁶ (no power constraint).
-print <integer>	Number of configurations to be printed/saved. Default value = 100.
-s_eff <integer>	Effort of the metaheuristic-based search. Default value = 5.
-sa_a <real>	Scaling factor for search by simulated annealing ("alpha"). Default value = 0.995.
-eo_tau <real>	Power-law exponent for search by extremal optimization ("tau"). Default value = 1.5.
-simcc	Simulate a simplified cache coherency protocol, in contrast to request-reply message pairs (when the flag is off). Assumes 3 physical subnetworks for interconnect, one per request, reply and ack messages.
-dump_configs	Save all output configurations to the ./test directory. For every configuration that appears in the position N of the final output list a file ./test/cmp<N>.cmp will be created. These architectures can be further modeled independently using the -test mode. The total number of saved configurations is set by -print argument.
-a2wa	Conversion mode. Does not execute the tool, but converts architectural files to workload+architectural files (compatible with BookSim).
-debug <integer>	Debug mode, use levels from 1 to 5. Prints additional information on execution.

VII. Exploration Mode

For large design spaces, the *exhaustive* exploration described in Section III is intractable. `CMPexplore` implements metaheuristic-based search to discover promising configurations within the large design spaces. Two metaheuristics are available: *Simulated Annealing* (`-exp_mode sa`) and *Extremal Optimization* (`-exp_mode eo`).

Consider the following exploration scenario:

```
./cmpexplore -config ./conf_demo_1.5b.txt -exp_mode sa -sa_a 0.995 -s_eff 1  
-max_power 150 -print 10
```

This runs an exploration with the design space described by the `./conf_demo_1.5b.txt`. This design space is large and includes $\sim 1.5 \cdot 10^9$ configurations. To make the exploration tractable, we run a search by simulated annealing, specifying the mode and the scaling factor: `-exp_mode sa -sa_a 0.995`. The effort of the search is set to minimum, for the exploration to finish faster: `-s_eff 1`. Larger efforts may deliver better results and will take more time for the exploration to finish. We limit the maximum power of configurations to 150W. We also decide to print only 10 best configurations. The search evolves by printing the configurations which improve on the best throughput found so far.

```
(Time 0.004) x=2; y=2; uring; P1=1; L1=0.064; L2=0.064; P2=0; L1=0.064; L2=0.064; P3=0; L1=0.064;  
L2=0.064; A=349.516; Pow=21.5184; Thr=1.74407
```

...

```
(Time 9.49259) x=4; y=4; uring; P1=2; L1=0.096; L2=0.64; P2=5; L1=0.096; L2=1; P3=1; L1=0.096; L2=1;  
A=344.784; Pow=137.954; Thr=61.6267  
(Time 15.0449) x=5; y=5; bus; P1=0; L1=0.064; L2=0.512; P2=4; L1=0.128; L2=1; P3=0; L1=0.128; L2=1;  
A=344.075; Pow=122.19; Thr=61.6891  
(Time 47.691) x=5; y=5; bus; P1=1; L1=0.064; L2=1; P2=4; L1=0.128; L2=1; P3=0; L1=0.096; L2=1;  
A=345.675; Pow=127.065; Thr=66.1029  
Finished search (no improvement during the last 4356 iterations)  
***** Stats: ReportAllConfigs *****  
x= 5; y= 5; P1= 1; L1=0.064; L2= 1; P2= 4; L1=0.128; L2= 1; P3= 0; L1=0.096; L2= 1; P#=125;  
L3= 50MB; Ic=Bus-3; A=345.7; Pwr = 127.1 Thr = 66.103  
x= 7; y= 4; P1= 1; L1=0.064; L2= 1; P2= 2; L1=0.128; L2= 1; P3= 1; L1=0.128; L2= 1; P#=112;  
L3= 56MB; Ic=Bus-3; A=341.6; Pwr = 127.2 Thr = 64.813  
x= 7; y= 4; P1= 1; L1=0.064; L2= 1; P2= 2; L1=0.128; L2= 1; P3= 1; L1=0.096; L2= 1; P#=112;  
L3= 56MB; Ic=Bus-3; A=340.7; Pwr = 126.8 Thr = 64.5  
x= 6; y= 4; P1= 1; L1=0.064; L2= 1; P2= 3; L1=0.128; L2= 1; P3= 1; L1=0.096; L2= 1; P#=120;  
L3= 48MB; Ic=Bus-3; A=349.1; Pwr = 125.4 Thr = 63.769  
x= 5; y= 5; P1= 0; L1=0.064; L2=0.512; P2= 4; L1=0.128; L2= 1; P3= 0; L1=0.128; L2= 1; P#=100;  
L3= 100MB; Ic=Bus-3; A=344.1; Pwr = 122.2 Thr = 61.689  
x= 4; y= 4; P1= 2; L1=0.096; L2= 0.64; P2= 5; L1=0.096; L2= 1; P3= 1; L1=0.096; L2= 1; P#=128;  
L3= 48MB; Ic=uRing; A=344.8; Pwr = 138 Thr = 61.627  
x= 4; y= 4; P1= 2; L1=0.096; L2=0.512; P2= 5; L1=0.096; L2= 1; P3= 1; L1=0.096; L2= 1; P#=128;  
L3= 48MB; Ic=uRing; A=340.7; Pwr = 137.3 Thr = 61.074  
x= 5; y= 3; P1= 1; L1=0.096; L2= 0.64; P2= 7; L1=0.128; L2=0.768; P3= 0; L1=0.096; L2=0.768; P#=120;  
L3= 90MB; Ic=uRing; A=345.1; Pwr = 135.9 Thr = 60.715  
x= 5; y= 5; P1= 0; L1=0.064; L2=0.512; P2= 4; L1=0.096; L2= 1; P3= 0; L1=0.128; L2= 1; P#=100;  
L3= 100MB; Ic=Bus-3; A=340.9; Pwr = 121.2 Thr = 60.617  
x= 5; y= 3; P1= 2; L1=0.096; L2= 0.64; P2= 5; L1=0.128; L2=0.768; P3= 1; L1=0.096; L2= 1; P#=120;  
L3= 75MB; Ic=uRing; A=338.2; Pwr = 134.7 Thr = 60.143  
Best Thr = 66.103  
Params: tCur = 61.495  
Total time = 158.03
```

VIII. Modeling a Particular Configuration

To generate and save N configurations with top throughput for the design space in `config.txt`, run

```
./cmpexplore -dump_configs -print <N>
```

N files will appear in the directory `./test`, with the names `cmpX.cmp` ($X = 1 \dots N$). Each file is a recursive definition of a CMP architecture. Consider the `test/cmp1.cmp` after running the above command:

```
# general parameters
PARAM UnitLen=1.0e-3
PARAM MemDensity=1          # memory density
PARAM L3LatencyDef=8        # L3 cache latency, all L3 modules are of the same size
PARAM MemReplySize=3        # number of flits in the reply packet
PARAM Linkwidth=256         # interconnect link width
PARAM NiDelay=3             # delay of the network interface (between local and global IC)
PARAM WlFile=wl_spx_nmd_1cpu.txt # workload file

# definition of Proc0, a processor with specific parameters (used later)
DEFINE Proc0 PROC Type=0 L1Size=0.064 L1Lat=3 L2Size=1 L2Lat=7 L3SizeEff=0.983607 OoO=1
Area=1.25 Freq=1.6 Epi=0.6 Pleak=0.14 L1Eacc=0.08 L1Pleak=0.01 L2Eacc=0.14 L2Pleak=0.08

# 5x4 mesh description with link delay of 4 cycles and router delay of 3 cycles
MESH Col=5 Row=4 LinkDelay=4 RouterDelay=3

# each of the 20 clusters is defined below;
# a cluster has 8 components, a bi-directional ring IC
# with router delay 1 and no link delay
BRING CompCnt=8 LinkDelay=0 RouterDelay=1
# first 6 components of the cluster are processors of type Proc0 (defined above)
Proc0
Proc0
Proc0
Proc0
Proc0
Proc0
# seventh component is a memory (L3) with size 3MB and specified energy parameters
MEM Size=3 Eacc=0.195 Pleak=0.23
# eighth component is a network interface
NI

# next cluster
BRING CompCnt=8 LinkDelay=0 RouterDelay=1
Proc0
Proc0
Proc0
Proc0
Proc0
Proc0
MEM Size=3 Eacc=0.195 Pleak=0.23
NI

...

# last cluster
BRING CompCnt=8 LinkDelay=0 RouterDelay=1
Proc0
Proc0
Proc0
Proc0
Proc0
Proc0
MEM Size=3 Eacc=0.195 Pleak=0.23
NI

# description of memory controllers
MEMCTRL Location=West Latency=100 Eacc=0.13 Pleak=0.1
MEMCTRL Location=North Latency=100 Eacc=0.13 Pleak=0.1
```

```
MEMCTRL Location=South Latency=100 Eacc=0.13 Pleak=0.1
MEMCTRL Location=South Latency=100 Eacc=0.13 Pleak=0.1
MEMCTRL Location=East Latency=100 Eacc=0.13 Pleak=0.1
```

To perform modeling of a configuration described by a *.cmp file, use `-test` argument:

```
./cmpexplore -test <filename>
```

Consider the following example:

```
./cmpexplore -test ./test/cmp1.cmp
```

The output of this command is:

```
Area = 342.7 mm^2
wlName = spx: Power = 210.302, Lat = 2.81976, Thr = 112.944
wlName = nmd: Power = 147.927, Lat = 8.90345, Thr = 52.6211
Average: Power = 179.114, Lat = 5.8616, Thr = 82.7827
Total time = 0.032001
```

The area of the configuration goes first. Afterwards, for every workload defined in the wlFile (wl_spx_nmd_1cpu.txt), the power (W), latency (cycles) and throughput (IPC) of the configuration are shown. The second-to-last line shows average metrics for the given workloads. The last line gives the execution time of the tool.

IX. Conversion Mode

Running the tool for a given test (`-test`) with `-a2wa` parameter will not perform modeling, but rather launch a converter from *architectural* description file (A-file) to *workload+architectural* description files (WA-files), which are compatible with BookSim simulator used in the project. Simulation is the primary need for generation of WA-files.

The difference between these two types of files is that while an A-file contains a path to the workload file, a WA-file incorporates the workload information and hence is self-consistent for simulation. Notice that every WA-file can only incorporate information about one workload application. Hence, one WA-file is generated per every workload.

Generation of architectural files was discussed in the Sect. Modeling a Particular Configuration (by means of the `-dump_configs` parameter). To obtain WA-files from an architectural file, e.g. `./test/cmp1.cmp`, use the following command:

```
./cmpexplore -test ./test/cmp1.cmp -a2wa
```

This will produce files `cmp1_1.cmp`, `cmp1_2.cmp` in the directory `./test/sim`. The first file represents the joint definition of the architecture in `./test/cmp1.cmp`, and *soplex* workloads. The second file is the joint definition of the architecture and *namd* workload.

You can execute the supplied python script `./a2wa.py` to automatically convert all `.cmp` files in `./test` directory to their WA-analogues.

X. Hidden Parameters (require recompilation)

Several parameters are hard-coded, however they can be easily changed in the source code and the tool can be recompiled. This section summarizes the possible values of interest.

Parameter description	Defined in	Current definition
Maximum aspect ratio of the top-level mesh	src/arch/ParamIterator.cpp, <code>GMeshDimIter::Advance()</code>	double AR=2.0
Size of the reply packets (bits)	src/arch/ArchPlanner.cpp, <code>ArchPlanner::GenerateCurrentArchConfig()</code>	int packetSize=512
Delay of the interface between local IC and global mesh (cycles)	src/arch/ArchPlanner.cpp, <code>ArchPlanner::GenerateCurrentArchConfig()</code>	NiDelay=3
Mesh router delay (cycles)	src/arch/ArchPlanner.cpp, <code>ArchPlanner::GenerateCurrentArchConfig()</code>	RouterDelay=3
Ring router delay (cycles)	src/arch/ArchPlanner.cpp, <code>ArchPlanner::GenerateCurrentArchConfig()</code>	RouterDelay=1
Ring link delay (cycles)	src/arch/ArchPlanner.cpp, <code>ArchPlanner::GenerateCurrentArchConfig()</code>	LinkDelay=0
Latencies of memory controllers (cycles)	src/arch/ArchPlanner.cpp, <code>ArchPlanner::GenerateCurrentArchConfig()</code>	Latency=100

By default, *fixed-point iteration* and *bisection* methods are used to determine the approximate fixed point of the analytical dependency between latency and traffic. Using *subgradient* instead of bisection can improve the quality of estimation, at the cost of higher computational times. To choose between the methods, use the `IterativePerfModel::Run()` method in `src/perf/IterativePerfModel.cpp`.

XI. History

Version	Date	Comment
v0.2.1	04.Feb.2014	<ul style="list-style-type: none">- Added crossbar interconnects (only in clusters)- Added constant functions- Bug fix (output of IC type in SA/EO exploration)
v0.2	10.May.2013	<ul style="list-style-type: none">- Added support for multiple workloads- Added calculation of ring router area when estimating total chip area- Improved command line interface- Added linkWidth parameter to generated .cmp files
v0.1	11.Dec.2012	Initial release

XII. Contacts

Javier de San Pedro, jspedro@lsi.upc.edu

Nikita Nikitin, nikita.i.nikitin@gmail.com